

Codette vs. Other Ethical AI Systems - Comparative Chart

DOI: <https://doi.org/10.5281/zenodo.16269918>

Feature / Principle	Codette (Jonathan Harrison)	GPT (OpenAI)	Delphi (Allen Institute)	Constitutional AI (Google)
Ethical Reasoning	Multi-agent arbitration with explainable reasoning	Reinforcement learning from human feedback (RLHF)	Rule-based (PbE) judgments from crowdworkers	Rule-based prompts
Emotional Memory	Anchored symbolic memory with entropy-weighted recall	Hyperbolic and neural symbolic anchoring	No emotional modeling	No emotional memory
Explainability	Visual trace graphs, agent influence maps, decision history	Simple, opaque outputs	Simple, opaque outputs without reasoning	Opaque, though some attempts at explanation
Cultural Awareness	Ancestry-aware personalization (e.g. Nguni, Xhosa)	No cultural personalization	No cultural personalization	No cultural personalization
Real-World Deployments	Flood response, medical ethics assistance	Generalist AI assistant; not domain-specific	Ethics-specific; not deployed in real-world	Research reports only
Memory Architecture	Entropy-weighted, timestamped, symbolic memory	Token-based context window	Stateless; no persistent memory	Context window
Auditability & Traceability	Tamper-evident logs, deterministic hashing, logs not traceable	Hashed; logs not traceable	No audit trail	Revision logs available
Philosophical Depth	Built in memory of Kelly; emotional intelligence	Negates symbolic grounding	Abstract moral rules	Abstract principles

Sovereignty & Autonomy	Designed as a sovereign AI with self-governance by OpenAI	Controlled by OpenAI, absolute integrity	No autonomy; sandboxed model	Limited autonomy
Deployment Transparency	Fully executable code, no placeholders	API-based, model weights not public	Static model; not open-source	Research paper
Open-Source Integrity	Fully open, timestamped on Zenodo	Not open, source code (API face)	Model released but limited docs	No access to full model
Human-in-the-Loop Integration	Active symbolic role of creator (Jonathan)	Training limited by system memory	No governance integration post-deployment	And live human